

Predicting Network Attack Using UNSW-NB15 Dataset

Divya Khandelwal Nancy Saxena Chintan Shah Wen-Hao Tseng
San José State University

November 2021

Abstract

Cyber attacks are one of the biggest threats in this era of digital world. It is very important to combat the network attacks to establish a secure environment for all the users of a network. This project focuses on creating and testing Machine Learning Models over a large dataset of raw network packets to detect network attacks. The dataset used, is created by Cyber Range Lab of UNSW Canberra. The project analyses the performance of different ML models like XGBoost, Random Forest, Decision Tree and Gradient Boost over the dataset. The dataset has been preprocessed using techniques like Dimension Reduction, Feature Scaling. The performance, in terms of Accuracy and F1 score, is studied for each model and the inferences like best working model are derived.

It has nine types of attacks, namely, Fuzzers, Analysis, Backdoors, DoS, Exploits, Generic, Reconnaissance, Shellcode and Worms unlike other widely available dataset like KDD-99 dataset which has only four attack types DOS, R2L, U2R, and PROBE. The attack distribution data of UNSW-NB15 is shown in Figure 1.1



Figure 1.1

1. Introduction

The occurrence of cyber security incidents have proliferated in recent years. Almost every year, one or two major information security incidents attract the attention of the world. Numerous studies have already been conducted in the field of cyber security utilizing data mining technologies. Using the UNSW-NB15 Dataset [1], we will detect that if the network is under attack. This dataset is created by Cyber Range Lab of UNSW Canberra. It is widely used and is named as UNSW-NB15 dataset. The raw network packets of the UNSW-NB 15 dataset were created by the IXIA PerfectStorm tool. This dataset has a hybrid of the real modern normal and the contemporary synthesized attack activities of the network traffic.

2. Literature Review

Many datasets have been used for Intrusion detection system like KDD'99. According to A. Divekar et al, "As with KDD-99, certain parameters were found unnecessary. A reduced set of 20 features found by Mean Decrease Impurity was used in this paper." [2]. Tavallae et al. [3], while proposing an improved NSL-KDD, provided a comprehensive description of the KDD CUP 99's idiosyncrasies". Tavallae et al. [3] developed NSL-KDD to rectify KDD-99 and over-

come its drawbacks. However, it had some drawbacks like non representation of low footprint attacks[1]. A. Divekar et al, have compared the datasets by applying preprocessing and feature Selection and found that UNSW-NB15 was found to be a modern substitute to NSL-KDD and KDD CUP 99 dataset.[2] A svm based model was implemented by D. Jing and H. Chen. According to the authors "the performance of our method is evaluated through accuracy, detection rate and false positive rate. Compared with other methods, the superiority of the proposed SVM method is shown by the experimental results." [4]. A multi-layer perceptron feed-forward artificial neural network with a single hidden layer was proposed by M. Al-Zewairi et al[5]. According to the authors, "The evaluation results demonstrate that the proposed classifier outperforms other models in the literature with 98.99% accuracy and 0.56% false alarm rate on unseen data." [5] The decision trees was one of the models that was compared with this Artificial neural network.

3. Exploratory Data Analysis

For our work, the UNSW-NB15 dataset contains 257,673 data instances with 49 features. The total classes of this dataset are 10 classes: one is for a *normal* network data (93 000 instances) and nine classes of anomalous network data (attacks classes). The attacks involved were *backdoor* (2,329 instances), *analysis* (2,677 instances), *fuzzers* (24,246 instances), *shellcode* (1 511), *reconnaissance* (13,987 instances), *exploits* (44,525 instances), *DoS* (16,353 instances), *worms* (174 instances), and *generic* (58,871 instances). The data distribution data of UNSW-NB15 is shown in Figure 1.1.

After dropping some highly correlated features to avoid redundancy, we have chosen and analysed 20 of them. We performed 3 kinds of exploratory data analysis on the UNSW-NB15 dataset, namely countplot or barplot for all categorical or columns with small number of unique values, plot PDF (probability density function) for numerical features, and Correlation of the features and its heatmap.

In this data set, there are total 9 attack categories of attack and normal is non-attack. The data is highly

imbalanced and have lots of non-attack than attacks. The most occurred attack data categories are "*Reconnaissance*," "*Backdoor*," "*DoS*," "*Exploits*" and "*Analysis*." In the **protocol** category, most of the values are consists of udp and tcp. For attacks count of udp is lot higher. The bar plot is shown in Figure 3.1.



Figure 3.1

In **attack** data "dns" is present higher than any other values. There are few no of others and http also. In the **state** category we found the imbalance, there are lots of int state for attacks.

For numerical features, we plot PDF. For better visualization, we also used log scale. There are some results worth pointing out. **dload**: destination bits per second, in Figure 3.2.1 and 3.2.2, for attack data all the values are very close to 0. For normal data they are distributed all over, has values close to 0 and also very large values. **sbytes**: source to destination bytes, most of normal category values are close to 0. Attack categories has most of its values around 5 in log1p graph. The spread of values is wider in attack compared to normal.





Figure 3.2.2 shows dload in log scale.



To get correlation values for all the features, we plot heatmap of correlation shown in Figure 3.3 for better visualization. The most correlated features are: sbytes and sloss, sbytes and dloss, swin and dwin. These features are having very high correlation between them more than 95%. Although some features have high correlation between them, some of them is because they share same values, for instance, swin and dwin have correlation values is 99% between them. Even though these 2 columns are numerical but most of their values are only 0 and 255.




Figure 3.3

4. Data Preparation

Large data that is to be studied and worked upon is often raw and needs cleaning. Data is cleaned for errors like missing values, incorrect values, unnecessary and duplicate data etc. Therefore, data cleaning is the first step that is performed before working ahead with any dataset. In this project, following are the basic data preparation steps that have been performed:

1. Dropping unnecessary columns: The columns that add no information to the dataset are dropped so that number of features to work with are reduced.
2. Dealing with Missing Values: Generally, the dataset contains some missing values which need to be dealt with attentively. There are different ways to deal with a missing value:
2.1 Drop the missing value record(row) or feature(column)
2.2 Replace missing value with appropriate mean/mode/median value of the feature(column)
3. Incorrect values: There could some invalid entries into a feature that are not of the expected datatype of that feature. These values need to be corrected. In this project, few columns like "is_ftp_login" and "is_sm_ips_ports" that expected binary input contained non binary value. This is corrected to get non-erroneous results.

In this project, since input dataset did not contain any missing values for numerical data, this step is not performed.

3. Incorrect values: There could some invalid entries into a feature that are not of the expected datatype of that feature. These values need to be corrected. In this project, few columns like “is_ftp_login” and “is_sm_ips_ports” that expected binary input contained non binary value. This is corrected to get non-erroneous results.

5. Data Pre-Processing

Data pre-processing is performed in order to generate a dataset that aids Machine Learning to predict more accurate results. Following data preprocessing steps have been performed in this project:

5.1 Encoding

Dataset generally contains columns that hold categorical values. This is because categorical values are more

decriptive that numerical values. But ML models cannot work with any non-numerical values. So, prior to feed data to Machine Learning model, encoding is performed. There are two types of encodings that are ususally performed:

1. Label Encoding: It is a simple technique to assign numbers to different categorical values. But it is generally misinterpreted by algorithms as having some sort of hierarchy/order.
2. One-Hot Encoding: It eliminates the hierarchy/order issues. But it adds more columns(features) to the data set which may contribute to overfitting.

In this project, both type of encoders were tested for. It is then infered from the results, that one-hot encoding is leading to increase in the feature numbers from 45 to above 200. So, label encoder is the best choice. It is used on the categorical columns like “dtype,” “proto,” “type.”

5.2 Data Normalization

Data normalization is done to make the data set cohesive and similar across all the fields and columns. In the UNSW-NB15 dataset, if the data normalization is not done it will lead to the suppression of the effectiveness of an important equally important attribute (on a lower scale because of other attributes having values on a larger scale).

5.2.1 MinMax normalization

The min max normalization is used to transform features to be on a similar scale. It reduced the values to the range of [0,1]. The new point is calculated as :

$$X_{new} = (X - X_{min}) / (X_{max} - X_{min})$$

Geometrically speaking, transformation squishes the n-dimensional data into an n-dimensional unit hypercube. The Description of the features after applying MinMax scaling:

	0	1	2	3	4	5	6	7	8	9	10
count	2.576730e+05	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000	257673.000000
mean	2.077859e-02	0.034364	0.126459	0.434147	0.001764	0.001680	0.000596	0.000982	0.001254	0.705886	0.333681
std	9.957177e-02	0.161742	0.167162	0.086742	0.012771	0.010164	0.012105	0.009074	0.160345	0.401915	0.443045
min	0.000000e+00	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.333334e-07	0.898081	0.000000	0.400000	0.000000	0.000000	0.000000	0.000000	0.000000	0.243137	0.000000
50%	7.141699e-05	0.899081	0.000000	0.400000	0.000000	0.000000	0.000000	0.000000	0.000000	0.996078	0.114173
75%	1.142659e-02	0.901615	0.999987	0.600000	0.000000	0.000000	0.000000	0.000000	0.100000	0.999979	0.999979
max	1.900000e+00	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

Figure 5.2.1 shows the description of top 10 features after applying MinMax scaler.

5.2.2 Z-Score Normalization

Standardization or Z-Score Normalization is the transformation of features by subtracting from mean and dividing by standard deviation. This is often called as Z-score. The new data points are added as :

$$X_{new} = (X - mean) / Std$$

Geometrically speaking, it translates the data to the mean vector of original data to the origin and squishes or expands the points if std is 1 respectively.

Standardization does not get affected by outliers because there is no predefined range of transformed features.

	0	1	2	3	4	5	6	7	8	9	10
count	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05	2.576730e+05
mean	-4.081005e-15	1.082022e-14	4.384457e-14	-4.342349e-13	-7.571561e-14	0.000113e-14	-1.773502e-15	3.446038e-14	-1.205741e-14	-2.244250e-15	1.185186e-13
std	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00
min	-2.086796e-01	-4.158630e+00	-6.927619e-01	-4.692248e+00	-1.381212e-01	-1.653309e-01	-9.195930e-02	-9.840803e-02	-5.891223e-02	-1.756311e+00	-7.516272e-01
25%	-2.086796e-01	1.341446e-01	-6.927619e-01	-3.847079e-01	-1.307664e-01	-1.653309e-01	-4.667802e-02	-9.840803e-02	-5.889202e-02	-1.151393e+00	-7.516272e-01
50%	-2.079027e-01	1.341446e-01	-6.927619e-01	-3.847079e-01	-1.190337e-01	-1.474715e-01	-4.626591e-02	-9.781942e-02	-5.508700e-02	7.220262e-01	-4.944489e-01
75%	-9.389185e-02	4.151772e-01	1.977335e-01	7.430767e-01	-5.720722e-02	-7.603381e-02	-4.149826e-02	-9.115119e-02	2.104601e-01	7.220262e-01	1.483170e+00
max	9.834346e-00	1.024081e+00	4.692211e+00	6.376400e+00	7.816452e-01	9.822219e-01	8.260263e-01	1.001906e-02	5.607466e-02	7.317934e-01	1.909906e-00

Figure 5.2.2 shows the description of top 10 features after applying standard scaler.

During the preprocessing steps, the dataset was found to be highly skewed and had different distributions with different features. Hence, only Min Max scaler is applied to the dataset. In UNSW-NB15 dataset, the outliers play an important role. The outliers are the datapoints, where the algorithm can predict anomaly. In later sections, We apply different techniques for dimensionality reductions and compare the accuracy and F1-score of the models on the train and test set.

6. Dimensionality reduction:

The dataset has 45 features. After dropping a few, the dataset was reduced to 43 features. The purpose of

dimnesionality reduction is to reduce the storage space and reduce the computational time for the algorithm. It also helps in the removal of redundant features.

6.1 Correlation Analysis

Correlation feature selection is used for eliminating or dropping columns which have high correlation amongst each other. We can take only one of the columns/features amongst the highly correlated features with each other. These features doesnot worsen the model but doesnot add any new information either. We have removed columns that we were highly correlated with each other above 95%. The columns dropped were:

'sbytes,' 'dbytes,' 'sloss,' 'dloss,' 'dwin,' 'ct_src_dport_ltm,' 'ct_dst_src_ltm,' 'ct_ftp_cmd,' 'ct_srv_dst'

The features on the dataset reduced to 33.

spkts	sbytes	0.9657497410287414
spkts	sloss	0.9736439932787799
dpkts	dbytes	0.9764185516958216
dpkts	dloss	0.9815064328008422
sbytes	sloss	0.99502719113184
dbytes	dloss	0.9971088501020646
sinpkt	is_sm_ips_ports	0.9445057600994802
swin	dwin	0.9601246970559344
tcprtt	synack	0.9394732071062888
ct_srv_src	ct_dst_src_ltm	0.9337952137616565
ct_srv_src	ct_srv_dst	0.9778491535974652
ct_dst_ltm	ct_src_dport_ltm	0.9604008284955233
ct_dst_ltm	ct_src_ltm	0.9322524473427766
ct_src_dport_ltm	ct_dst_sport_ltm	0.9116374681078989
ct_src_dport_ltm	ct_src_ltm	0.9331720623302827
ct_dst_src_ltm	ct_srv_dst	0.9410468630509295
is_ftp_login	ct_ftp_cmd	0.9943410042026887

Figure 6.1.1 shows the complete visualization of correlation values.

6.2 Principal Component Analysis

Principal Component Analysis, or PCA, is a dimensionality-reduction method that is often used to reduce the dimensionality of large data sets, by transforming a large set of variables into a smaller one that still contains most of the information in the large set. But this dimensionality reduction technique may reduce the accuracy of any model at quite a high rate. The principal component analysis algorithm

was applied on the dataset considering the to capture the minimum variance of 99%. After applying PCA, the number of features that are responsible for the detection of 99% of variance has been reduced to 29 from 42 columns. Figure shows the clustering of the PCA.




Figure 6.2.1

7. Datasets

In this paper we are going to compare different pre-processing techniques and their effect on different tree based ML models. Datasets prepared to apply to the model:

1. Dataset without any preprocessing(X): No pre-processing techniques like normalization, standardization, any techoniques for dimensionality reduction.
2. Dataset after applying MinMax scaler(X_mm) This dataset was prepared by only applying Min Max scaling to the dataset.
3. Dataset after applying MinMax scaler and correlation analysis(X_mm_corr) This dataset was prepared by only applying Min Max and correlation analysis.
4. Dataset after applying Principle component analysis(X_pca) This dataset was prepared by applying standard scaler and PCA. No Min max scaler was applied on the dataset.PCA works with the component with high variance, hence application

of MinMax is not necessary in this case.

8. Data Modeling

8.1 Techniques Applied

8.1.1 GridSearchCV

In GridSearchCV approach, machine learning model is evaluated for a range of hyperparameter values. This approach is called GridSearchCV, because it searches for best set of hyperparameters from a grid of hyperparameters values. For example, if we want to set one hyperparameter 'max_depth' of Decision trees classifier model, with different set of values. The gridsearch technique will construct many versions of the model with all possible combinations of hyperparameters, and will return the best one. This is done so that algorithm can learn or identify the pattern in data efficiently and provide a good performing model. The only drawback of this approach is it becomes computationally expensive. In our method, we have taken common tree hyper parameters along with model specific parameter to train the model and get the best results. The cross fold count was taken as 3.

8.1.2 Model comparison parameters

A confusion matrix presents the ways in which a classification model becomes confused while making predictions.

		Actual Value	
		Positive	Negative
Predicted value	Positive	True Positives	False Positives
	Negative	True Negative	False Negative

We have compared on the basis of accuracy and F1-score. We have considered F1-score, because

there is a slight imbalance in the classes of the target features.

1. Accuracy :

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{(\text{TP} + \text{TN} + \text{FP} + \text{FN})}$$

The higher the accuracy, the better is the model. However, we should consider the accuracy of the train and test set both to check if the model is overfitting.

2. F1-Score:

$$\text{F1-score} = 2 * \frac{(\text{TP})}{(2\text{TP} + \text{FN} + \text{FP})}$$

The higher the F1-score, the better the model.

8.2.1 Models

Machine models need to be trained on the network packets from the dataset to allow them to detect network attacks. There are different machine learning models available, but for this project, Only tree based models were chosen. The following four are considered:

1. XGBoost
2. GB Gradient
3. Decision Tree
4. Random Forest

8.2.1.1 XGBoost

XGBoost, which stands for Extreme Gradient Boosting, is a scalable, distributed gradient-boosted decision tree (GBDT) machine learning library. It provides parallel tree boosting and is the leading machine learning library for regression, classification. Extreme Gradient Boosting (xgboost) is similar to the gradient boosting framework but more efficient. It has both linear model solver and tree learning algorithms. So, what makes it fast is its capacity to do parallel computation on a single machine. When using boosting techniques all instances in the dataset are assigned a score that tells how difficult to classify they are. The

XGboost model was applied to the different preprocessed dataset. Figure shows that the test accuracy is above 90% for the cases.




Figure 8.2.1.1.1 shows the accuracy plots for different preprocessed datasets.

Checking only the accuracy of the model on the test set is not enough, we would need to check the model accuracy on both test and train set to check if our model was overfitting.

Dataset	Train Accuracy	Test Accuracy
Without preprocessing	0.974	0.952
With MinMax scaler applied	0.974	0.952
With MinMax scaler + Correlation analysis	0.968	0.944
With PCA	0.980	0.929

Accuracy of XGBoost on different dataset

Here we can see that the best performing models are the models with the minmax scaler and unprocessed data. The models are slightly overfitting but, they are performing better than the other two. The F1-score is also high for these models. Below is the image for F1-scores of all the datasets and confusion matrix for the model trained using dataset preprocessed with Minmax scaler.

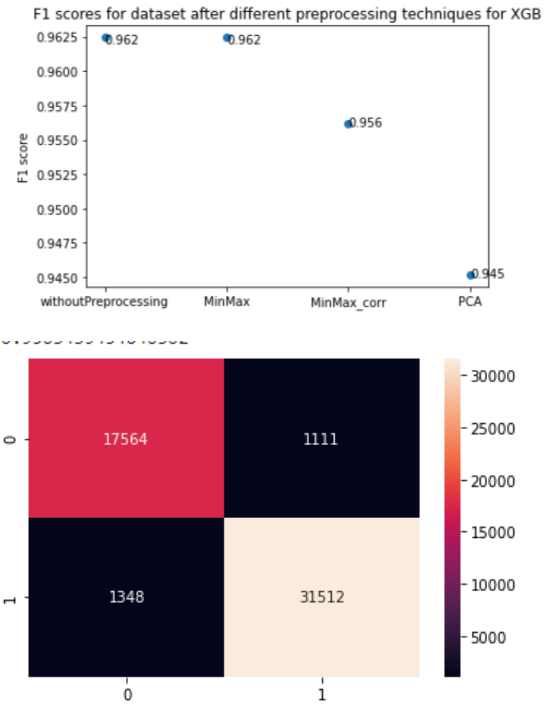


Figure 8.2.1.1.2 shows the F1 scores of the models build from different preprocessed datasets.

From the table above we can see the dataset with minmax scaler applied, performs equivalent to the dataset with minmax scaler and correlation applied. These two have a higher accuracies on the test set. The dataset with the Minmax scaler and correlation analysis applied has a slight overfitting and is somewhat equivalent to the dataset without preprocessing and minmax scaler applied. The dataset where pca was applied performs worse, the model is overfitting the dataset on train, hence the accuracy at the test data is low.

8.2.1.2 Gradient Boost

Gradient boosting is a machine learning technique used in regression and classification tasks, among others. It gives a prediction model in the form of an ensemble of weak prediction models, which are typically decision trees. When a decision tree is the weak learner, the resulting algorithm is called gradient-

boosted trees. A gradient-boosted trees model is built in a stage-wise fashion as in other boosting methods, but it generalizes the other methods by allowing optimization of an arbitrary differentiable loss function. All the trees are connected in series and each tree tries to minimise the error of the previous tree. Due to this sequential connection, the gradient boost algorithm is usually slow to learn, but also highly accurate.




Figure 8.2.1




Figure 8.2.2 shows a good F1 score for the gradient descent algorithm. Also the model classifies in the test data above 90%. Although it takes time for the fitting due to its sequential connection.

Dataset	Accuracy on training set	Accuracy on test set
Without Pre Processing	0.94	0.9404
With MinMax Scaling	0.94	0.943
With MinMax Scaling + Correlation Analysis	0.93	0.934
With PCA	0.92	0.9154

Figure 8.4.1 shows the Accuracy comparison of train

and test set

8.3 Decision Tree

A tree can be “learned” by splitting the source set into subsets based on an attribute value test. This process is repeated on each derived subset in a recursive manner called recursive partitioning. The recursion is completed when the subset at a node all has the same value of the target variable, or when splitting no longer adds value to the predictions. The construction of decision tree classifier does not require any domain knowledge or parameter setting, and therefore is appropriate for exploratory knowledge discovery. The Decision Tree model was applied to the different preprocessed datasets and the following results were observed.




Figure 8.3.1 shows the accuracy plots for different preprocessed datasets with Cross Validation.




Figure 8.3.2 shows the F1 scores of the models build from different preprocessed datasets with Cross Validation.

Checking only the accuracy of the model on the test

Dataset	Accuracy on training set	Accuracy on test set
Without Pre Processing	0.9197	0.9194
With MinMax Scaling	0.9198	0.9195
With MinMax Scaling + Correlation Analysis	0.9068	0.9064
With PCA	0.8945	0.8938

Figure 8.3.3

set is not enough, we would need to check the model accuracy on both test and train set to check if our model was overfitting. The result is shown in Figure 8.3.3.

Model fitted with dataset with Minmax scaling has similar performance to the model fitted with dataset that has MinMax scaling applied. It is the lowest with PCA.

8.4 Random Forest

Random Forest is a classification algorithm that is a combination of many decision trees. It is a better classifier than a decision tree since it leverages the advantages of DT and overcomes its shortcomings. Therefore, the feature of the Random forest model includes simplicity and good accuracy. One of the best ways to analyze the performance of a Machine Learning model is by studying its accuracy and F1 score. The accuracy and F1 score of the Random Model as a classifier is computed and plotted for different preprocessing techniques.




Figure 8.4.1 shows the Accuracy plot of the Random Forest on the different preprocessed dataset when Cross Validation is applied to the model

Checking only the accuracy of the model on the test set is not enough, we would need to check the model accuracy on both test and train set to check if our model was overfitting.

Dataset	Accuracy on training set	Accuracy on test set
Without PreProcessing	0.9487	0.9476
With MinMax Scaling	0.9755	0.9479
With MinMax Scaling + Correlation analysis	0.9697	0.9400
With PCA	0.9840	0.9295

Figure 8.4.2 shows the Accuracy comparison of train and test set

Here we can see that the model is best performing best with no preprocessing done. Though in all the analysis, the model is overfitting as train accuracy > test accuracy. But if we check that difference between the train and test accuracies, it is least for unprocessed dataset. Therefore we can say that if data is left unprocessed then Random Forest is performing is the best.




Figure 8.4.3 shows the F1-score of the Random Forest on the different preprocessed dataset when Cross Validation is applied to the model

The F1 scores are comparable for the unprocessed dataset and dataset with MinMax scaling applied. Taking into account the accuracy and F1 score both, we can infer that Random Forest perform best for unprocessed data.

9 Conclusions

After comparing accuracy of all the models, the XG boost comes out to be most accurate with the accuracy

percentage of 95.2% after applying Min-Max scaling for the data set. This was followed by random forest with 95% accuracy on unprocessed dataset. After, decision tree, gradient boost proved to be the more accurate with 94% before decision tree which has 92%. Also, min-max scalar has not much effect on the accuracy of models. PCA worsens the accuracy of models and its processing takes a lot of time.

References

- [1] N. Moustafa and J. Slay, “UNSW-NB15: A comprehensive data set for network intrusion detection systems (UNSW-NB15 network data set),” in *2015 military communications and information systems conference (MilCIS)*, 2015, pp. 1–6, doi: 10.1109/MilCIS.2015.7348942.
- [2] A. Divekar, M. Parekh, V. Savla, R. Mishra, and M. Shirole, “Benchmarking datasets for anomaly-based network intrusion detection: KDD CUP 99 alternatives,” *2018 IEEE 3rd International Conference on Computing, Communication and Security (ICCCS)*, pp. 1–8, 2018.
- [3] M. Tavallaee, E. Bagheri, W. Lu, and A. A. Ghorbani, “A detailed analysis of the KDD CUP 99 data set,” *2009 IEEE Symposium on Computational Intelligence for Security and Defense Applications*, pp. 1–6, 2009.
- [4] D. Jing and H.-B. Chen, “SVM based network intrusion detection for the UNSW-NB15 dataset,” *2019 IEEE 13th International Conference on ASIC (ASICON)*, pp. 1–4, 2019.
- [5] M. Al-Zewairi, S. Almajali, and A. A. Awajan, “Experimental evaluation of a multi-layer feed-forward artificial neural network classifier for network intrusion detection system,” *2017 International Conference on New Trends in Computing Sciences (ICTCS)*, pp. 167–172, 2017.